

```
title: "Predictive Analysis and Model Evaluation"  
author: "Adedeji Ogundipe"  
  
output:  
  
  word_document: default  
  
  pdf_document: default  
  
---  
  
```{r setup, include=FALSE}  
  
knitr::opts_chunk$set(echo = TRUE)  
  
...  
  
```{r}  
  
rm(list = ls())  
  
library(readr)  
  
test_data <- read.csv("C:/Users/User/Desktop/PaPA/CW-test.csv")  
  
train_data <- read.csv("C:/Users/User/Desktop/PaPA/CW-train.csv")  
  
...  
  
Test_data Variables
```

```
```{r}  
  
names(test_data)  
  
...  
  
[1] "age"        "job"         "marital"      "education"    "default"     "balance"  
[7] "housing"     "loan"        "contact"     "day"          "month"       "duration"  
[13] "campaign"    "pdays"       "previous"    "poutcome"    "y"  
...  
X
```

Train\_data Variables

```
```{r}  
  
names(train_data)
```

```
[1] "age"      "job"       "marital"    "education" "default"   "balance"
[7] "housing"   "loan"      "contact"    "day"       "month"     "duration"
[13] "campaign"  "pdays"     "previous"   "poutcome"  "y"
```

Number of subscribers based on train\_data

```
```{r}
```

```
table(train_data$y)
```

0	1
27920	3728

As 0 = "no" and 1 = "y", there are 3,728 subscribers and 27,920 non-subscribers.

Indicating "job" and other variables as factors in R

```
```{r}
```

```
glm(y ~ as.factor(job) + as.factor(housing)+ as.factor(education), family='binomial',
data = train_data)
```

```
...
```

```
Call: glm(formula = y ~ as.factor(job) + as.factor(housing) + as.factor(education),
family = "binomial", data = train_data)
```

```
Coefficients:
```

	(Intercept)	as.factor(job)blue-collar
as.factor(job)entrepreneur	-1.82069	-0.39093
as.factor(job)management	-0.53131	-0.49765
as.factor(job)self-employed	-0.22831	0.52241
as.factor(job)student	-0.34587	-0.33802
as.factor(job)unemployed	0.76392	-0.20853
as.factor(housing)yes	0.03381	as.factor(education)secondary
as.factor(education)tertiary	-0.74823	0.23250
	0.60389	as.factor(education)unknown
		0.29681

```
Degrees of Freedom: 31647 Total (i.e. Null); 31632 Residual
Null Deviance: 22950
Residual Deviance: 21930          AIC: 21960
```

Obtaining the standard errors and p-values of the train\_data

```
```{r}
```

```

Model1logit = glm(y ~ as.factor(job) + as.factor(housing)+ as.factor(education),
family='binomial', data = train_data)

summary(Model1logit)

```

```

Call:
glm(formula = y ~ as.factor(job) + as.factor(housing) + as.factor(education),
family = "binomial", data = train_data)

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-0.9921 -0.5540 -0.4333 -0.3559  2.5116 

Coefficients:
              Estimate Std. Error z value Pr(>|z|)    
(Intercept) -1.82069  0.08030 -22.675 < 2e-16 ***
as.factor(job)blue-collar -0.39093  0.07209 -5.423 5.87e-08 ***
as.factor(job)entrepreneur -0.53131  0.12389 -4.289 1.80e-05 *** 
as.factor(job)housemaid   -0.49765  0.13365 -3.723 0.000197 *** 
as.factor(job)management   -0.22831  0.07185 -3.178 0.001485 ** 
as.factor(job)retired      0.52241  0.08193  6.376 1.81e-10 *** 
as.factor(job)self-employed -0.34587  0.11072 -3.124 0.001785 ** 
as.factor(job)services      -0.33802  0.08374 -4.037 5.42e-05 *** 
as.factor(job)student       0.76392  0.10296  7.419 1.18e-13 *** 
as.factor(job)technician    -0.20853  0.06763 -3.083 0.002048 ** 
as.factor(job)unemployed    0.03381  0.10789  0.313 0.754005    
as.factor(job)unknown       -0.54155  0.23571 -2.298 0.021587 *  
as.factor(housing)yes       -0.74823  0.03771 -19.844 < 2e-16 *** 
as.factor(education)secondary 0.23250  0.06295  3.693 0.000221 *** 
as.factor(education)tertiary  0.60389  0.07203  8.384 < 2e-16 *** 
as.factor(education)unknown  0.29681  0.10267  2.891 0.003840 ** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 22945  on 31647  degrees of freedom
Residual deviance: 21926  on 31632  degrees of freedom
AIC: 21958

Number of Fisher Scoring iterations: 5

```

```{r}

```

predict_model1 <- predict(Model1_train, train_data, type = "response")

```

```

Building a logistic model to predict using a probability cut-off of 0.5

```{r}

```

install.packages("ggplot2", repos = 'http://cran.us.r-project.org')

```

```

installed.packages("lattice")

```

```

install.packages("caret")
```
```
```{r}
library(ggplot2)
library(lattice)
library(caret)
```
```
```{r}
predict_model1_class <- ifelse(predict_model1 >0.5, 1, 0)
train_data$predict_model1_class = predict_model1_class
confusionMatrix(as.factor(train_data$predict_model1_class), as.factor(train_data$y))
```
```
Warning: Levels are not in the same order for reference and data. Refactoring data to match.
Confusion Matrix and Statistics
Prediction Reference
      0      1
0 27920 3728
1      0      0
Accuracy : 0.8822
95% CI  : (0.8786, 0.8857)
No Information Rate : 0.8822
P-Value [Acc > NIR] : 0.5044
Kappa : 0
McNemar's Test P-Value : <2e-16
Sensitivity : 1.0000
Specificity : 0.0000
Pos Pred Value : 0.8822
Neg Pred Value : 0.0000
Prevalence : 0.8822
Detection Rate : 0.8822
Detection Prevalence : 1.0000
Balanced Accuracy : 0.5000
'Positive' Class : 0

```

The report above suggests that job is not a sufficiently good predictor or determinant for subscription.

Building a model for "y" with more variables

```

```{r}
Model2_train = glm(y ~ age + as.factor(job) + as.factor(marital) + balance +
as.factor(housing) , family='binomial', data = train_data )

```

```

predict_model2 <- predict(Model2_train, train_data, type = "response")

predict_model2_class <- ifelse(predict_model2 > 0.5, 1, 0)

train_data$predict_model2_class = predict_model2_class

confusionMatrix(as.factor(train_data$predict_model2_class), as.factor(train_data$y))

...

```

### Confusion Matrix and Statistics

		Reference	
		Prediction	
		0	1
Prediction		0	27916 3727
		1	4 1
Accuracy : 0.8821			
95% CI : (0.8785, 0.8856)			
No Information Rate : 0.8822			
P-Value [Acc > NIR] : 0.5252			
Kappa : 2e-04			
McNemar's Test P-Value : <2e-16			
Sensitivity : 0.9998567			
Specificity : 0.0002682			
Pos Pred Value : 0.8822172			
Neg Pred Value : 0.2000000			
Prevalence : 0.8822042			
Detection Rate : 0.8820779			
Detection Prevalence : 0.9998420			
Balanced Accuracy : 0.5000625			
'Positive' Class : 0			

Even when more variables are considered, the prediction seems to be more at variance with reality. This may be an indictment on the model.

By adjusting the cut-off to 0.25,

```
```{r}
```

```

Model2_train = glm(y ~ age + as.factor(job) + as.factor(marital) + balance +
as.factor(housing) , family='binomial', data = train_data )

predict_model2 <- predict(Model2_train, train_data, type = "response")

```

```

predict_model2_class <- ifelse(predict_model2 > 0.25, 1, 0)

train_data$predict_model2_class = predict_model2_class

confusionMatrix(as.factor(train_data$predict_model2_class), as.factor(train_data$y))

...

```

#### Confusion Matrix and Statistics

		Reference	
		0	1
Prediction	0	27170	3366
	1	750	362

Accuracy : 0.8699  
 95% CI : (0.8662, 0.8736)  
 No Information Rate : 0.8822  
 P-Value [Acc > NIR] : 1  
 Kappa : 0.1009  
 Mcnemar's Test P-Value : <2e-16  
 Sensitivity : 0.9731  
 Specificity : 0.0971  
 Pos Pred Value : 0.8898  
 Neg Pred Value : 0.3255  
 Prevalence : 0.8822  
 Detection Rate : 0.8585  
 Detection Prevalence : 0.9649  
 Balanced Accuracy : 0.5351  
 'Positive' Class : 0


From the foregoing, the true positive is about 32.55% whilst the false positive takes about 64.45 of the total number of subscribers predicted. On the other hand, the prediction for non-subscribers has 88.98 true negative and 11.02 false negative. Consequently this model seems more reliable for predicting the number of potential non-subscribers than for predicting potential subscribers.

#### **Decision trees model**

As an alternative to the logistic model, the decision-trees model can be tried and its accuracy assessed accordingly.

```
```{r}
install.packages("rpart")
install.packages("rpart.plot")
```
```
```{r}
library(rpart)
library(rpart.plot)
```
```
```{r}
tree <- rpart(formula = y ~ job +age + marital+ balance+ duration + campaign + contact
+ loan + default ,
data = train_data,
# minsplit =5,
# minbucket =10,
# control = list(minbucket=10, maxdepth=4),
method = 'class', cp =0.004
)
```
```
```{r}
rpart.plot(tree)
```
```

```



```{r}

```

PredictCART_train = predict(tree, data = train_data, type = "class")
train_data$PredictCART_train = PredictCART_train
confusionMatrix(train_data$PredictCART_train, as.factor(train_data$y))
```
```

```

## Confusion Matrix and Statistics

|            |       | Reference |
|------------|-------|-----------|
| Prediction | 0     | 1         |
| 0          | 27499 | 3091      |
| 1          | 421   | 637       |

Accuracy : 0.889  
95% CI : (0.8855, 0.8925)  
No Information Rate : 0.8822  
P-Value [Acc > NIR] : 7.604e-05  
  
Kappa : 0.2259  
  
McNemar's Test P-Value : < 2.2e-16  
  
Sensitivity : 0.9849  
Specificity : 0.1709  
Pos Pred Value : 0.8990  
Neg Pred Value : 0.6021  
Prevalence : 0.8822  
Detection Rate : 0.8689  
Detection Prevalence : 0.9666  
Balanced Accuracy : 0.5779  
  
'Positive' Class : 0

Based on the decision tree model, the true positive records 60.21% as against the 39.79% recorded by false positive for the number of subscribers predicted. Contrarily, the prediction for non-subscribers has 89.90% true negative and 10.10% for false negative. This shows that the model still predicts the number of potential non-subscribers fairly better than that of the subscribers although there's an improvement when compared to the logistic model.

## KNN Model

Using the KNN model, a new data set is created from the train\_data set with 5000 observations and renamed "small\_train\_data". Another one is also created from the test\_data set and renamed small\_test\_data.

```{r}

```

small_train_data = train_data[1:1000,]

small_test_data = train_data[1:1000,]

...

```{r}

knn.fit <- train(y ~ age + balance , data = small_train_data,
method = "knn",
# tuneLength = 17,
preProcess=c("center", "scale"))

```

```

```

### k-Nearest Neighbors

```

1000 samples
  2 predictor
  2 classes: '0', '1'

Pre-processing: centered (2), scaled (2)
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 1000, 1000, 1000, 1000, 1000, 1000, ...
Resampling results across tuning parameters:

  k  Accuracy   Kappa
  5  0.9725655 -0.007521545
  7  0.9762180 -0.002869700
  9  0.9783961  0.0000000000
```

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was k = 9.

### Predicting with KNN using the small train\_data set

```

```{r}

predictions = predict(knn.fit, small_train_data)

confusionMatrix(predictions, small_train_data$y)
```

...

### Confusion Matrix and Statistics

Prediction	Reference	
	0	1
0	980	20
1	0	0

Accuracy : 0.98  
95% CI : (0.9693, 0.9877)  
No Information Rate : 0.98  
P-Value [Acc > NIR] : 0.5591

Kappa : 0

McNemar's Test P-Value : 2.152e-05

Sensitivity : 1.00  
Specificity : 0.00  
Pos Pred Value : 0.98  
Neg Pred Value : NaN  
Prevalence : 0.98  
Detection Rate : 0.98  
Detection Prevalence : 1.00  
Balanced Accuracy : 0.50

'Positive' Class : 0

Predicting with KNN using the original train\_data set

```{r}

```
knn.fit1 <- train(y ~ age + balance , data = train_data,  
method = "knn",  
# tuneLength = 17,  
preProcess=c("center", "scale"))
```

...

```{r}

```
predictions = predict(knn.fit1, train_data)  
confusionMatrix(predictions, train_data$y)
```

...

```

Confusion Matrix and Statistics

              Reference
Prediction      0      1
      0  27735   3419
      1    185    309

          Accuracy : 0.8861
          95% CI  : (0.8826, 0.8896)
No Information Rate : 0.8822
P-Value [Acc > NIR] : 0.01532

          Kappa : 0.1222

McNemar's Test P-Value : < 2e-16

          Sensitivity : 0.99337
          Specificity  : 0.08289
Pos Pred Value : 0.89025
Neg Pred Value : 0.62551
          Prevalence  : 0.88220
          Detection Rate: 0.87636
Detection Prevalence: 0.98439
Balanced Accuracy : 0.53813

'Positive' Class : 0

```

Using the original train\_data, KNN predicts the number of potential subscribers at 62.63% which appears marginally better than the position recorded with the decision-trees model.

## Random Forest Model

```

```{r}
install.packages("randomForest")
```
```
```{r}
library(randomForest)
```
```

```

```

```{r}

random_forest_class = randomForest(y ~.,
data = train_data, #train_data data set
importance = T)
```
```
```{r}

p1 <- predict(random_forest_class, train_data)

confusionMatrix(p1, train_data$y)
```
```

```

### Confusion Matrix and Statistics

		Reference	
		0	1
Prediction	0	27920	78
	1	0	3650

Accuracy : 0.9975  
 95% CI : (0.9969, 0.9981)  
 No Information Rate : 0.8822  
 P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.988


Mcnemar's Test P-Value : < 2.2e-16

	Sensitivity	Specificity
Pos Pred Value	0.9972	0.9791
Neg Pred Value	1.0000	0.9972
Prevalence	0.8822	0.8822
Detection Rate	0.8822	0.8847
Detection Prevalence	0.8847	0.8847
Balanced Accuracy	0.9895	0.9895

'Positive' Class : 0


This model records 99.51% accuracy and 100% true positive of the prediction for potential subscribers. So far, it appears the best of the models tested.

```
```{r}  
plot(random_forest_class)  
```
```



```
```{r}  
varImpPlot(random_forest_class)  
```
```

random\_forest\_class



Applying the models to the test data

### Logistic model

```
```{r}
predict_model2 <- predict(Model2_train, test_data, type = "response")
predict_model2_class <- ifelse(predict_model2 > 0.25, 1, 0)
test_data$predict_model2_class = predict_model2_class
confusionMatrix(as.factor(test_data$predict_model2_class), as.factor(test_data$y))
````
```

```

Confusion Matrix and Statistics

            Reference
Prediction      0      1
      0  11678  1404
      1    324   157

        Accuracy : 0.8726
        95% CI   : (0.8669, 0.8782)
  No Information Rate : 0.8849
P-Value [Acc > NIR] : 1

        Kappa : 0.1053

McNemar's Test P-Value : <2e-16

        Sensitivity : 0.9730
        Specificity  : 0.1006
  Pos Pred Value : 0.8927
  Neg Pred Value : 0.3264
        Prevalence  : 0.8849
        Detection Rate : 0.8610
  Detection Prevalence : 0.9645
  Balanced Accuracy : 0.5368

'Positive' Class : 0

```

Model accuracy: 87.26%

### Decision trees Model

```

```{r}

tree1 <- rpart(formula = y ~ job +age + marital+ balance+ duration + campaign +
contact + loan + default ,
data = test_data,
# minsplit =5,
# minbucket =10,
# control = list(minbucket=10, maxdepth=4),
method = 'class', cp =0.004
)

```

```

```

```{r}

PredictCART_test = predict(tree1, test_data, type = "class")

test_data$PredictCART_test = PredictCART_test

confusionMatrix(test_data$PredictCART_test, as.factor(test_data$y))

```

```

### Confusion Matrix and Statistics

|            |       | Reference |  |
|------------|-------|-----------|--|
| Prediction | 0     | 1         |  |
| 0          | 11807 | 1196      |  |
| 1          | 195   | 365       |  |

Accuracy : 0.8974  
 95% CI : (0.8922, 0.9025)  
 No Information Rate : 0.8849  
 P-Value [Acc > NIR] : 1.807e-06  
 Kappa : 0.3017  
 McNemar's Test P-Value : < 2.2e-16  
 Sensitivity : 0.9838  
 Specificity : 0.2338  
 Pos Pred Value : 0.9080  
 Neg Pred Value : 0.6518  
 Prevalence : 0.8849  
 Detection Rate : 0.8705  
 Detection Prevalence : 0.9587  
 Balanced Accuracy : 0.6088  
 'Positive' Class : 0

Model accuracy: 89.74%

### KNN Model

```

```{r}

predictions = predict(knn.fit, small_test_data)

confusionMatrix(predictions, small_test_data$y)

```

Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	980	20
1	0	0

Accuracy : 0.98  
95% CI : (0.9693, 0.9877)

No Information Rate : 0.98  
P-Value [Acc > NIR] : 0.5591

Kappa : 0

McNemar's Test P-Value : 2.152e-05

Sensitivity : 1.00  
Specificity : 0.00  
Pos Pred Value : 0.98  
Neg Pred Value : NaN  
Prevalence : 0.98  
Detection Rate : 0.98  
Detection Prevalence : 1.00  
Balanced Accuracy : 0.50

'Positive' Class : 0

Model accuracy: 98.00%

## Random forest

```{r}

```
random_forest_class1 = randomForest(y ~.,  
data = test_data, importance = T)
```

...

```
Warning: The response has five or fewer unique values. Are you sure you want to do regression?
```

```
```{r}
p2 <- predict(random_forest_class1, test_data)
confusionMatrix(p2, test_data$y)
```
```

```

The decision tree is hereby adopted as the choice model for the assignment

Assuming all the variables are included in the model

```
```{r}
tree2 <- rpart(formula = y ~ age + job + marital + education + default + balance +
housing + loan + contact + day + month + duration + campaign + pdays + previous +
poutcome ,
data = train_data,
# minsplit =5,
# minbucket =10,
# control = list(minbucket=10, maxdepth=4),
method = 'class', cp =0.004
)
```
```

```

Identifying the importance of each variable in the model (Goman, 2014a)


```
```{r}
tree2$variable.importance
```
```

```

	duration	poutcome	month	contact	day
1037.9492602	588.2552789	295.5053183	47.8173947	15.3466371	
housing	job	pdays	age	balance	
8.6261065	8.5196710	5.5441366	5.2497339	4.4469251	
campaign	previous	education	default		
2.0859253	1.8792353	0.4759936	0.3829057		

From the above, "duration" has the highest sum of the goodness of split measures and can therefore be considered as the most significant of the variables (Goman, 2014a; Therneau et al, 2022).

### Plotting the decision tree for tree2



### Checking the cross-validation results

```
```{r}
printcp(tree2)
```

```

```

Classification tree:
rpart(formula = y ~ age + job + marital + education + default +
      balance + housing + loan + contact + day + month + duration +
      campaign + pdays + previous + poutcome, data = train_data,
      method = "class", cp = 0.004)

```

Variables actually used in tree construction:  
[1] contact duration month poutcome

Root node error: 3728/31648 = 0.1178

n= 31648


|   | CP        | nsplit | rel error | xerror  | xstd     |
|---|-----------|--------|-----------|---------|----------|
| 1 | 0.0346924 | 0      | 1.00000   | 1.00000 | 0.015383 |
| 2 | 0.0249464 | 3      | 0.89592   | 0.93079 | 0.014910 |
| 3 | 0.0217275 | 4      | 0.87098   | 0.88948 | 0.014615 |
| 4 | 0.0073766 | 5      | 0.84925   | 0.86347 | 0.014424 |
| 5 | 0.0054319 | 7      | 0.83450   | 0.85354 | 0.014350 |
| 6 | 0.0042918 | 11     | 0.81277   | 0.84523 | 0.014288 |
| 7 | 0.0040000 | 13     | 0.80418   | 0.84469 | 0.014284 |

From the above, the best xerror is 0.84281 with xstd of 0.014270; hence, we look for the smallest tree with xerror less than 0.85708 i.e., (0.84281+0.014270). Therefore, the tree with cp=0.0040000 is selected and the tree will be pruned with a cp slightly greater than 0.004, say, 0.0041.

```

```
tree2 <- prune(tree2, cp = 0.0041)
rpart.plot(tree2)
```

```



Using the model to predict with test\_data

tree3 is created from tree1 with an adjustment from cp=0.004 to cp=0.0041

```
tree3 <- rpart(formula = y ~ job +age + marital+ balance+ duration + campaign + contact + loan + default ,
```

```
data = test_data,
```

```
# minsplit =5,
```

```
# minbucket =10,
```

```
# control = list(minbucket=10, maxdepth=4),
```

```
method = 'class', cp =0.0041
```

```
)
```

```
...
```

```
```{r}
```

```
PredictCART_test = predict(tree3, test_data, type = "class")
```

```
test_data$PredictCART_test = PredictCART_test
```

```
confusionMatrix(test_data$PredictCART_test, as.factor(test_data$y))
```

### Confusion Matrix and Statistics

		Reference	
Prediction		0	1
0	0	11807	1196
	1	195	365

Accuracy : 0.8974

95% CI : (0.8922, 0.9025)

No Information Rate : 0.8849

P-Value [Acc > NIR] : 1.807e-06

Kappa : 0.3017

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9838

Specificity : 0.2338

Pos Pred Value : 0.9080

Neg Pred Value : 0.6518

Prevalence : 0.8849

Detection Rate : 0.8705

Detection Prevalence : 0.9587

Balanced Accuracy : 0.6088

'Positive' Class : 0